

Exploring the Apache Spark Structured Streaming API for Processing Streaming Data

EXPLORING SOURCES AND SINKS



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Understanding Spark DataFrames

Data sources, data sinks, data transformations

Writing data to the console sink, file sink

Customizing write logic using the foreach sink and the foreachBatch sink

Prerequisites and Course Outline

Prerequisites

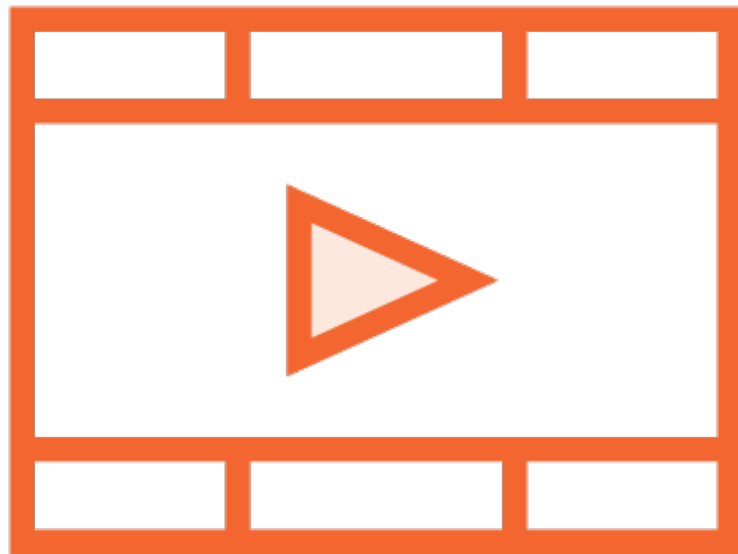


Comfortable programming in Python

**Some exposure to Apache Spark
and DataFrames**

**Some exposure to working with
streaming data using Apache Spark**

Prerequisite Courses



**Conceptualizing the Processing
Model for Apache Spark
Structured Streaming**

Course Outline



Exploring Sources and Sinks

Processing Streaming DataFrames

**Performing Windowing Operations
on Streams**

Working with Streaming Joins

**Managing and Monitoring
Streaming Queries**

Streaming DataFrames in Spark 2.x

DataFrame: Data in Rows and Columns

DATE	OPEN	...	PRICE
2016-12-01	772	...	779
2016-11-01	758	...	747
2006-01-01	302	...	309

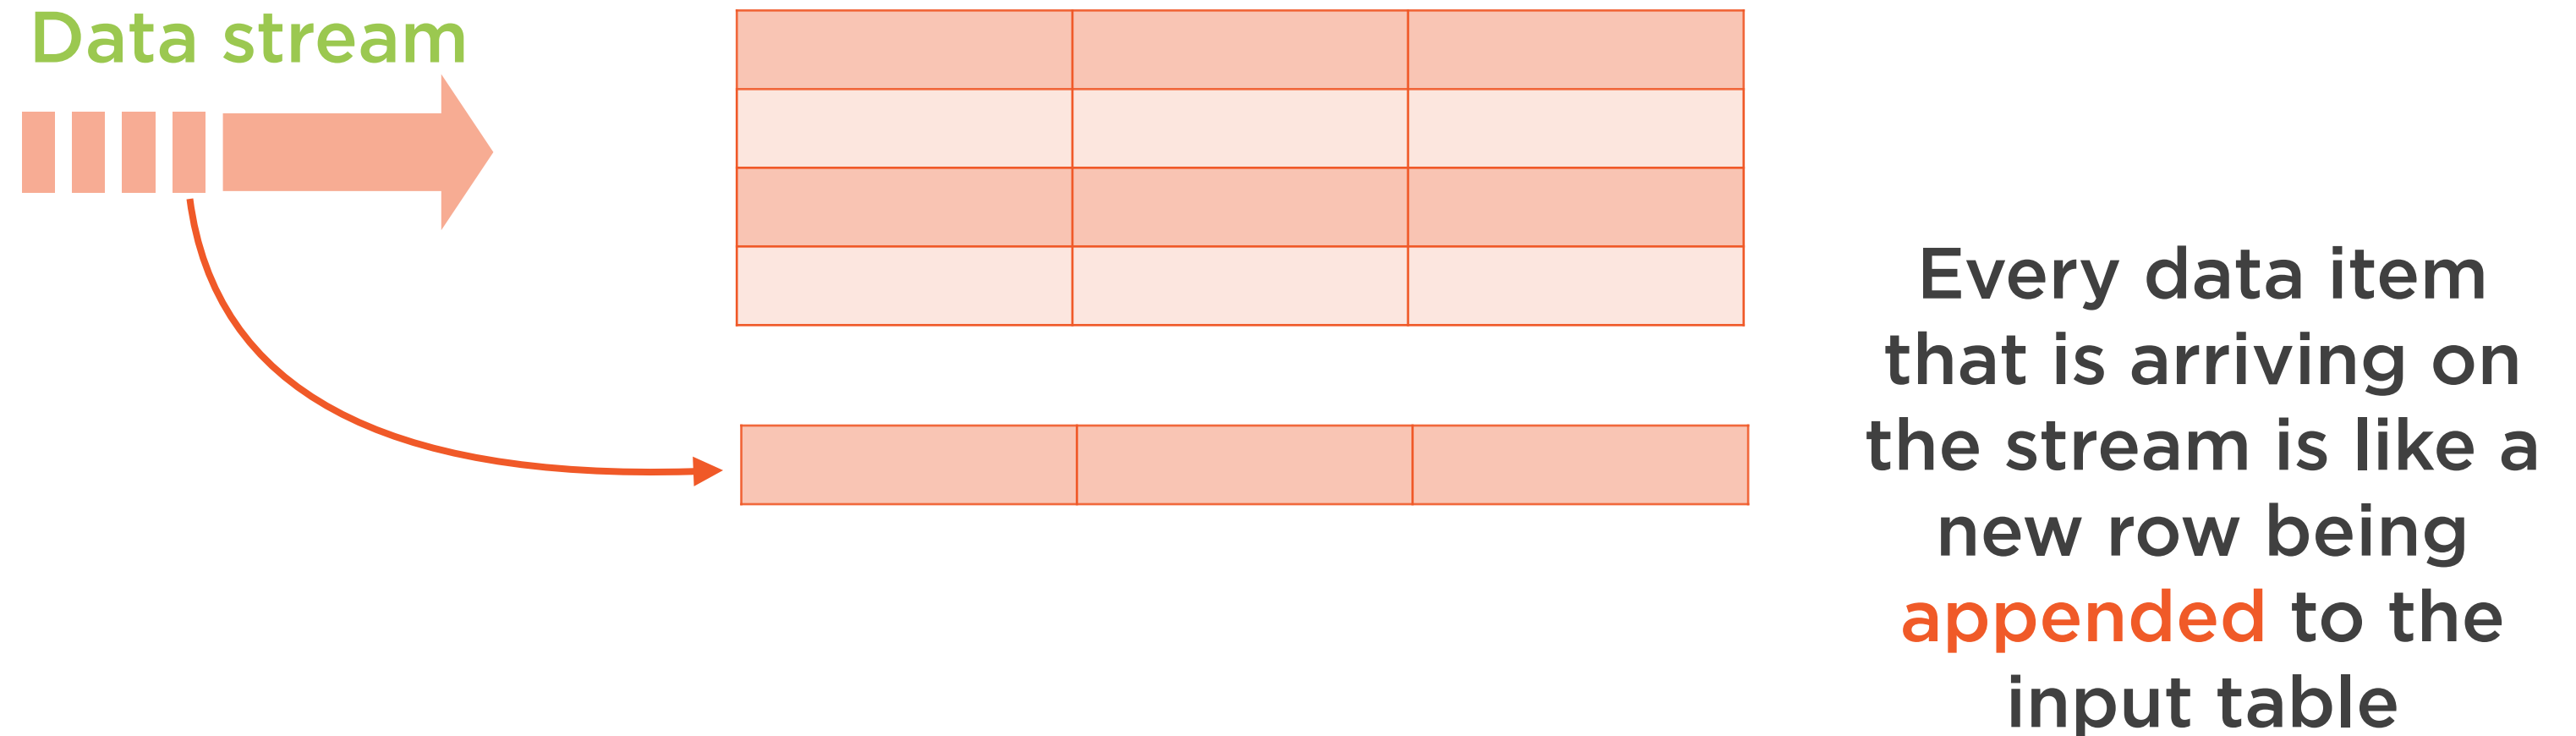
Streaming Data Spark 2.x

Data stream



**Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table**


Streaming Data Spark 2.x



Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream




Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



--	--	--

--	--	--


--	--	--

Every data item that is arriving on the stream is like a new row being **appended** to the input table

Data stream as an unbounded input table

Streaming Data Spark 2.x

Data stream



--	--	--

--	--	--

--	--	--


--	--	--

Every data item
that is arriving on
the stream is like a
new row being
appended to the
input table

Data stream as an unbounded input table

Batch is Simply A Prefix of Stream

Data stream



--	--	--

--	--	--


--	--	--

--	--	--

**In other words, the
input table (batch)
is simply a prefix of
the stream**

Batch is Simply A Prefix of Stream

Data stream



--	--	--

--	--	--

--	--	--

--	--	--

**All operations that
can be performed on
dataFrames can be
performed on the
stream**

Structured Streaming treats a live data stream as a table that is being **continuously** appended

Prefix Integrity

Running job on continuous data yields same result as running job on batch data (where the batch is a prefix or snapshot of continuous data)


Prefix Integrity

Running job on continuous data yields **same result** as running job on batch data (where the batch is a prefix or snapshot of continuous data)

Prefix Integrity

Running job on continuous data yields same result as running job on batch data (where the batch is a prefix or snapshot of continuous data)

Structured Streaming



What

A high-level API that
takes burden off user



How

Micro-batch
processing with
exactly-once fault-
tolerance



Why

Code virtually
identical for batch
and streaming

Structured Streaming



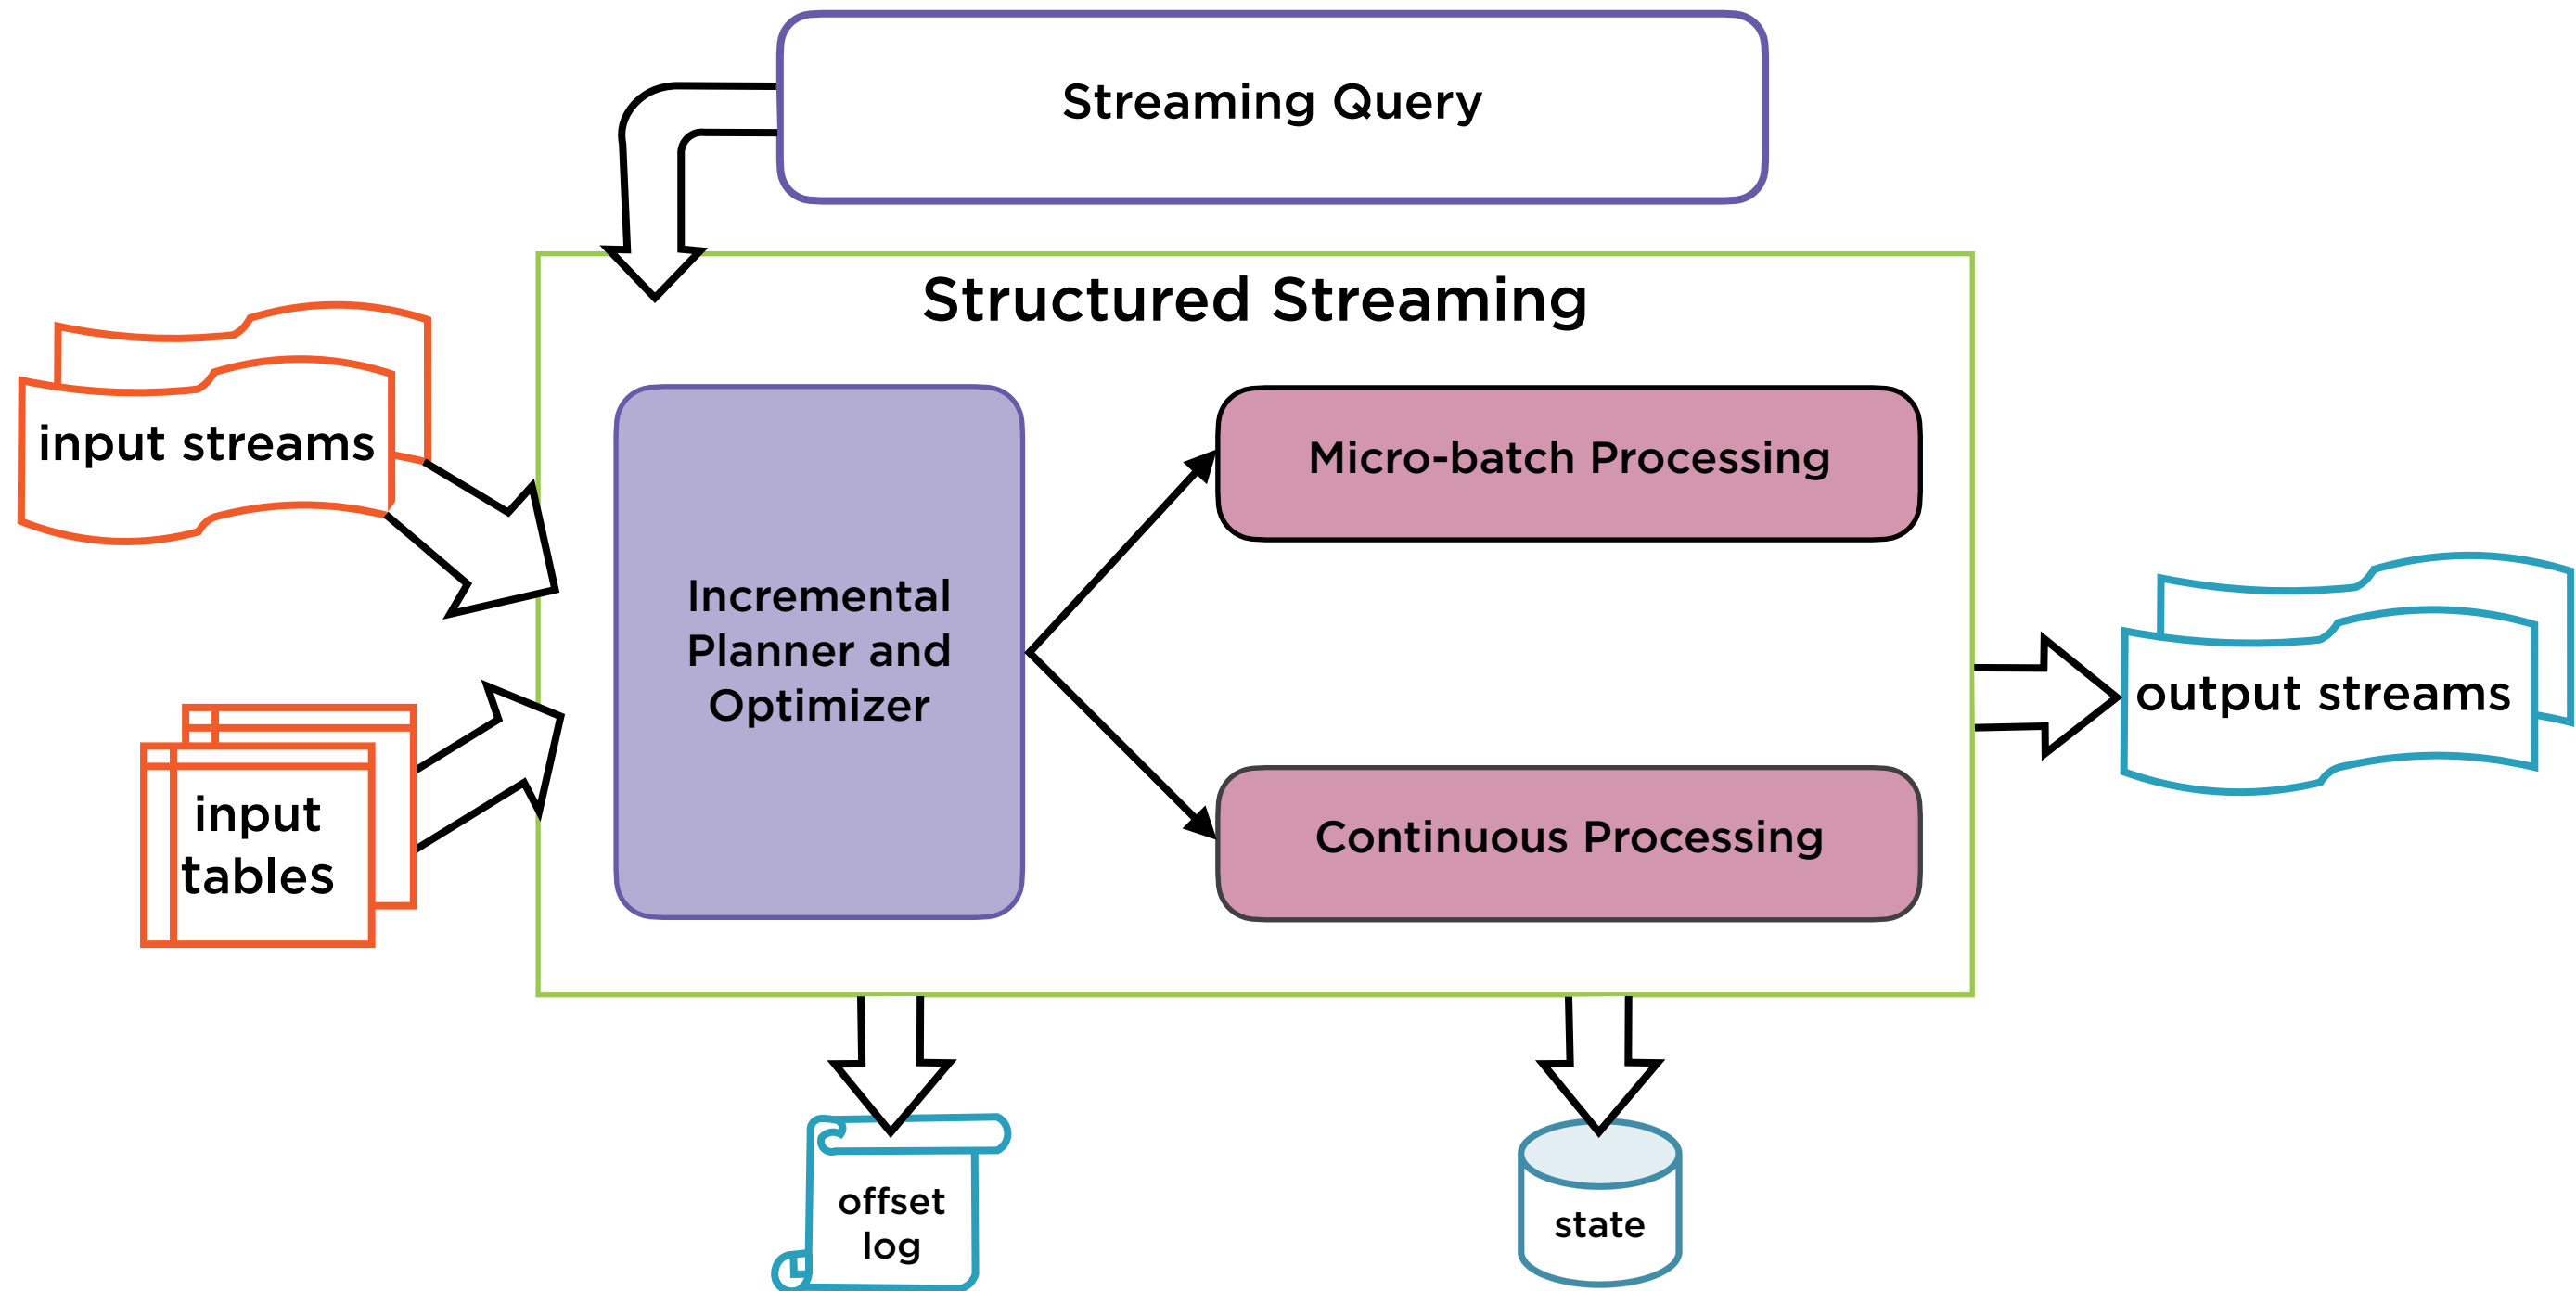
A stream is not a stream

- It is simply an unbounded dataset

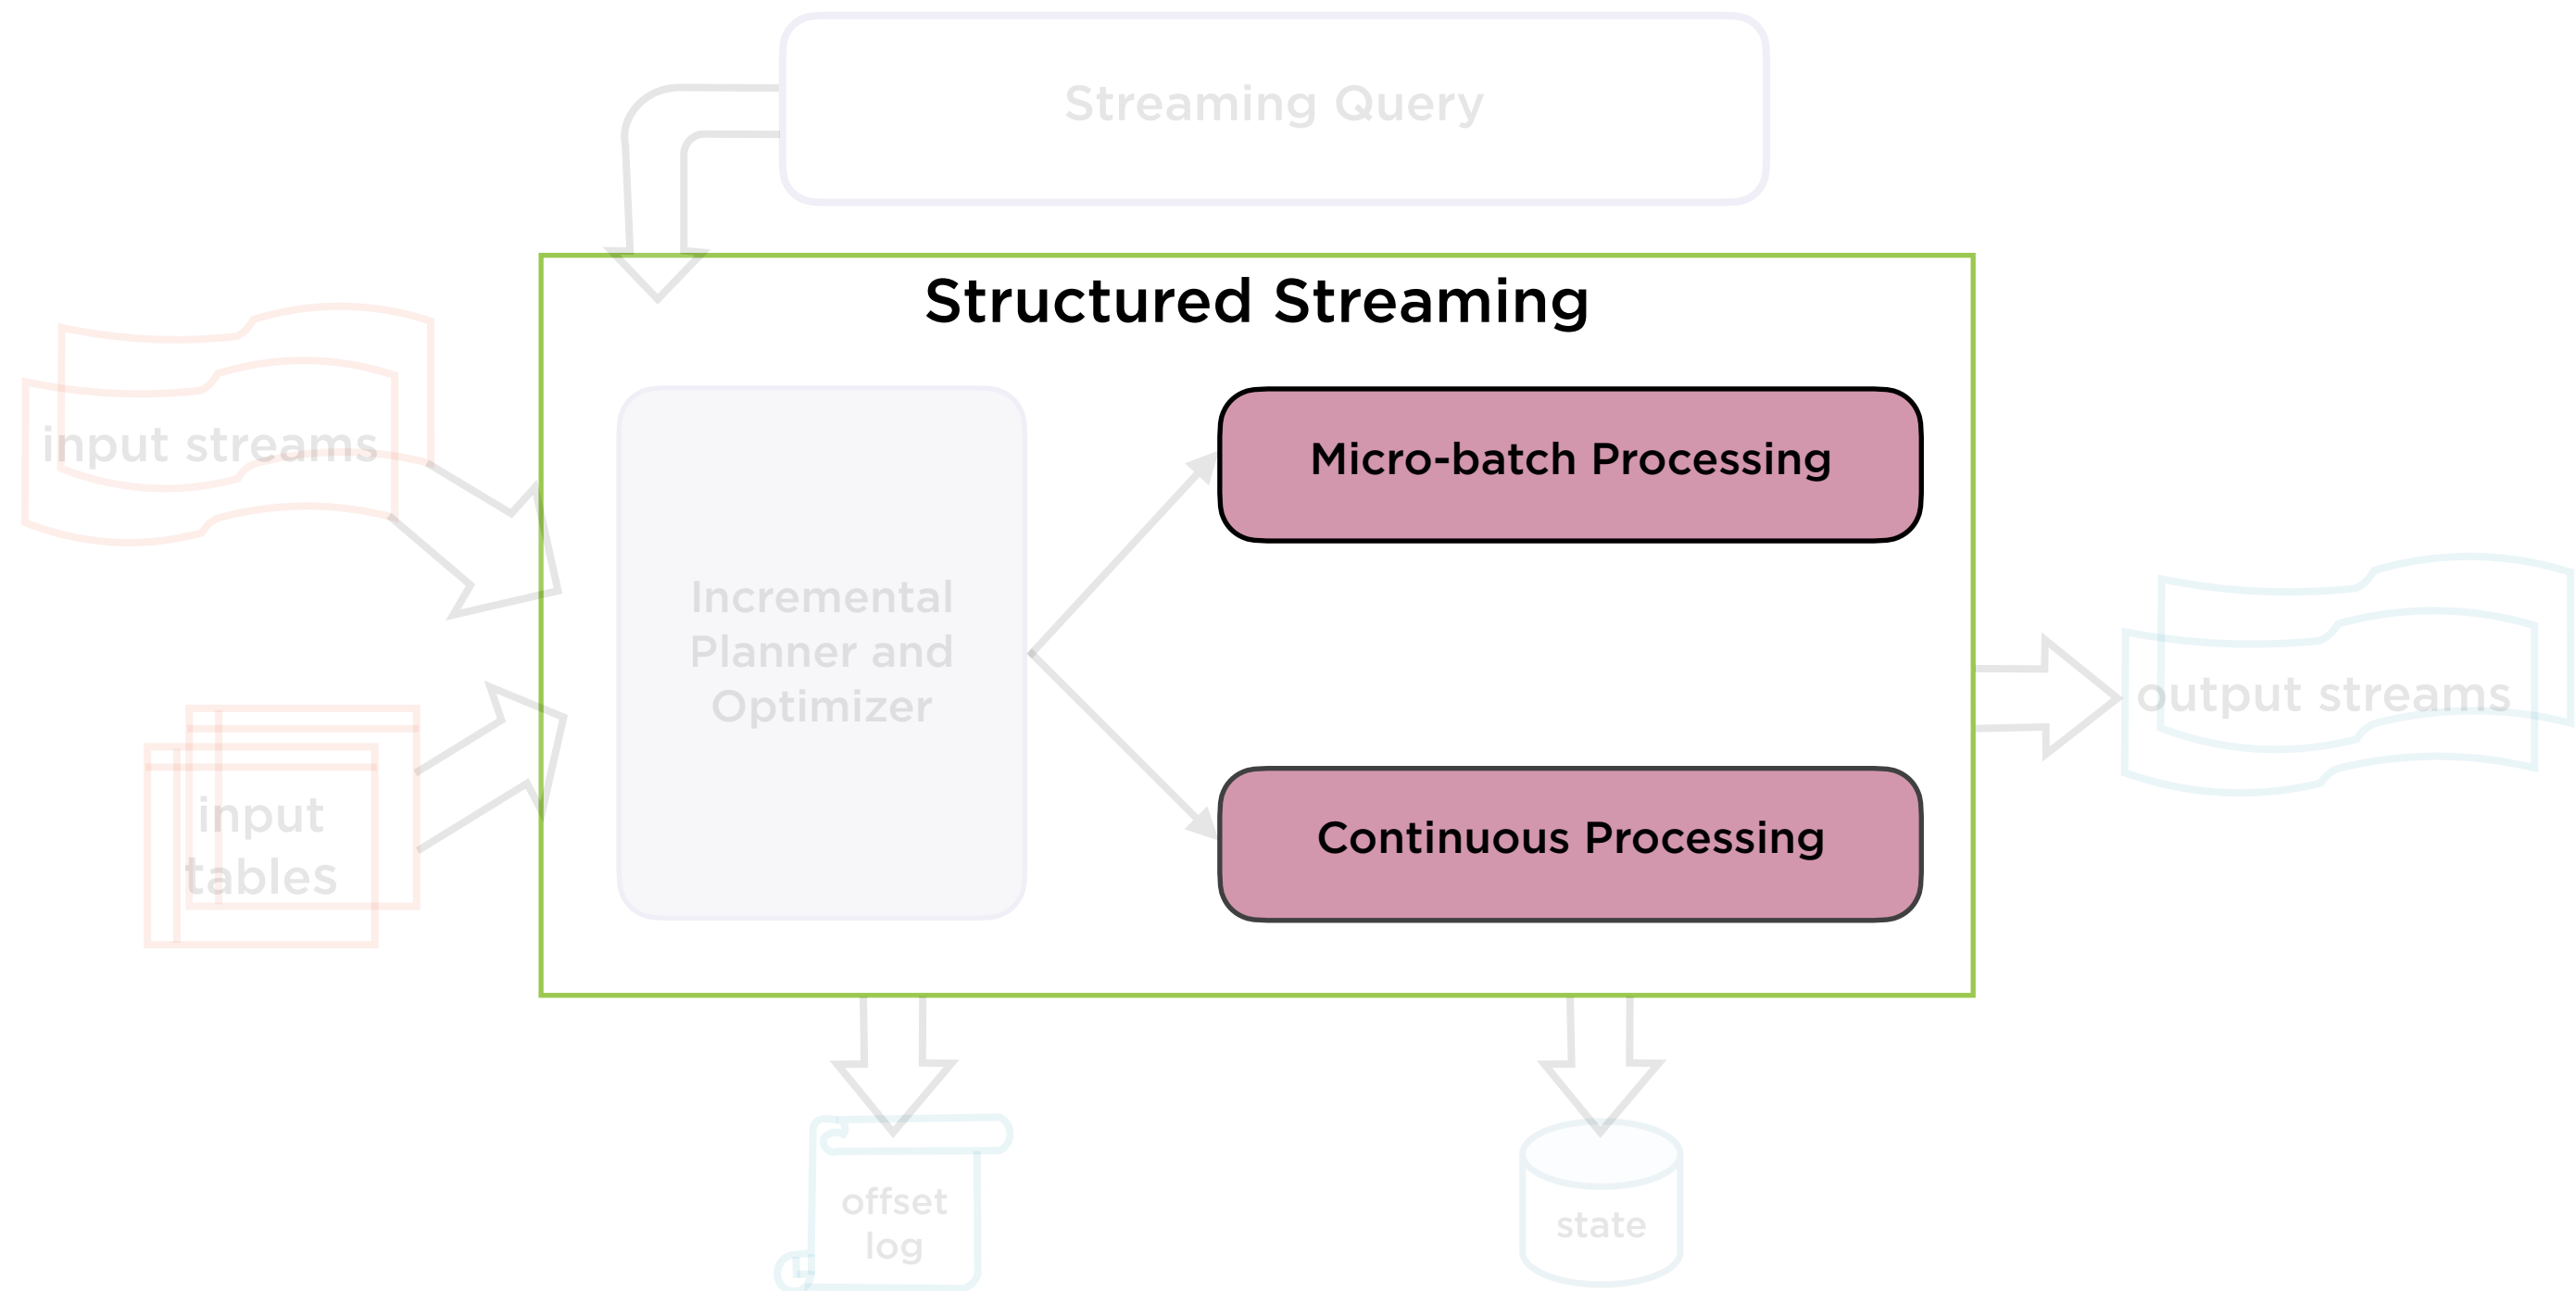
End-to-end application

- Unify batch and streaming pipelines
- With same queries for both

Micro-batch and Continuous Processing



Micro-batch and Continuous Processing



Micro-batch Processing in Spark



Default stream processing mode in Spark

Data streams processed as a series of batch jobs

End-to-end latencies as low as 100ms

Exactly-once fault-tolerance guarantees

Continuous Processing in Spark



Experimental stream processing mode in Spark


Data streams processed using long-running tasks

End-to-end latencies a few milliseconds

At least-once fault-tolerance guarantees

Built-in Sources and Sinks

Transformations on Streaming Data





Data Source


File source

Kafka source

Socket source

Rate source

File Source



Reads files written in a directory as a stream of data

Files processed in the order of modification time (or reverse order)

Files atomically placed in source directory

CSV, JSON, ORC, Parquet files


Kafka Source



**Open source software to handle
real-time data feeds**

**Compatible with Kafka broker
version 0.10.0 or higher**


Socket Source



Read UTF-8 text data from a socket connection

Test source, does not provide end-to-end fault tolerance guarantees

Rate Source



Generates data at the specified number of rows per second

Each row contains timestamp and value

Can be used for benchmarking streaming application code

Fault Tolerance of Sources




Different types of sources and sinks differ in their fault tolerance

For a source to be fault tolerant, it must be possible to replay data

- Then, can replay all data after checkpoint to restore system state

Fault Tolerance of Sources



Fault tolerant sources

- File sources
- Kafka sources
- Rate sources

Socket sources are not fault tolerant



Data Sink

File sink


Kafka sink

Foreach sink

Console sink

Memory sink

File Sink



Writes files to the specified output directory

Offers exactly-once fault tolerance semantics

Supports the append output mode

Kafka Sink

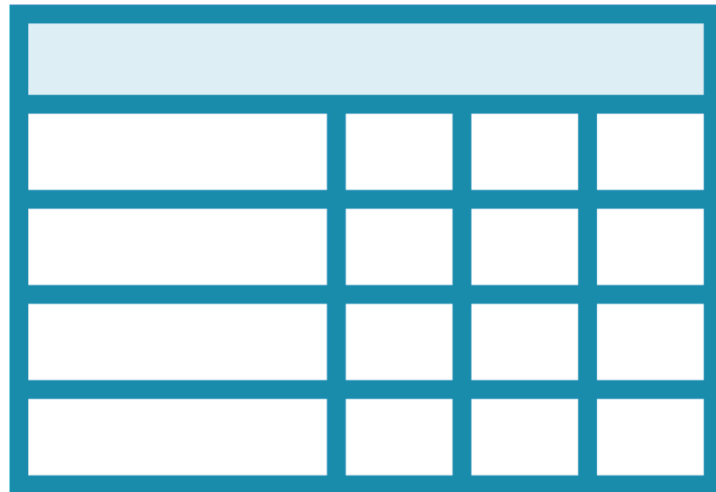


Writes output to one or more topics in Kafka

Supports at-least-once fault tolerance semantics

Supports the append, complete, and update output mode

Foreach Sink



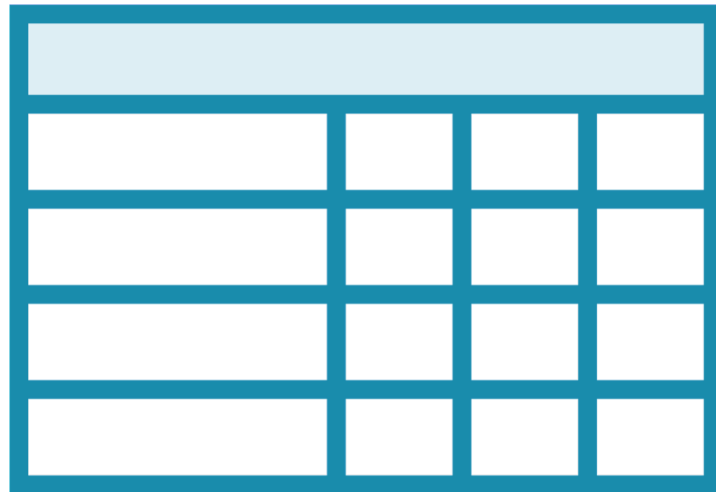
Custom write logic for results

Perform write operations for each row

Supports at-least-once fault tolerance semantics

Supports the append, complete, and update output mode

ForeachBatch Sink




Custom write logic for results

Perform write operations for each micro-batch processed

Fault tolerance semantics depends on the implementation

Supports the append, complete, and update output mode

Console Sink




Prints results out to the terminal window

Used for testing and debugging purposes

No fault tolerance guarantees

Supports the append, complete, and update output mode

Memory Sink



Output stored in memory in a table format

Used for testing and debugging purposes

No fault tolerance guarantees

Supports the append and complete output mode

Demo

Exploring the environment set up

Demo

Writing results to the console sink

Demo

**Writing results to the file sink using
CSV and JSON files**

Demo

Writing results using custom logic with the foreach sink

Demo

**Writing results using custom logic with
the foreachBatch sink**

Summary

Understanding Spark DataFrames

Data sources, data sinks, data transformations

Writing data to the console sink, file sink

Customizing write logic using the foreach sink and the foreachBatch sink

Up Next:

Processing Streaming DataFrames
