

# Performing Windowing Operations on Streams

---



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Windowing operations on streams**

**Sliding and tumbling windows**

**Event time vs. processing time**

**Aggregation operations on windows**

**UDFs on streaming data**

# Stateless and Stateful Transformations

---

# Transformations



## Stateless

Transformations which are applied on a single stream entity



## Stateful

Transformations which accumulate across multiple stream entities







# Stateless Transformations



Each entity is operated on standalone

Speed exceeded? **Alert** triggered





1





2





4




# Window Transformations




**Accumulate information across a window in a stream**


# Streaming Data



# Streaming Data




# Streaming Data



A window is a **subset** of a stream based on

- Time interval
- Count of entities
- Interval between entities

# Streaming Data



Transformations can be applied on all entities **within** a window

- sum, min, max, average



# Tumbling, Sliding, and Global Windows

---

# Types of Windows

**Tumbling Window**

**Sliding Window**

**Count Window**

**Session Window**

**Global Window**

# Types of Windows

**Tumbling Window**

**Sliding Window**

**Count Window**

**Session Window**


**Global Window**

# Types of Windows



**A stream of data**

# Tumbling Window




**Fixed window size**

**Non-overlapping time**


**Number of entities differ within a window**

# Tumbling Window




**The window tumbles over the data, in a non-overlapping manner**

# Tumbling Window



Apply the `sum()` operation on each window

# Sliding Window




Fixed window size

**Overlapping** time - sliding interval

Number of entities differ within a window



# Sliding Window

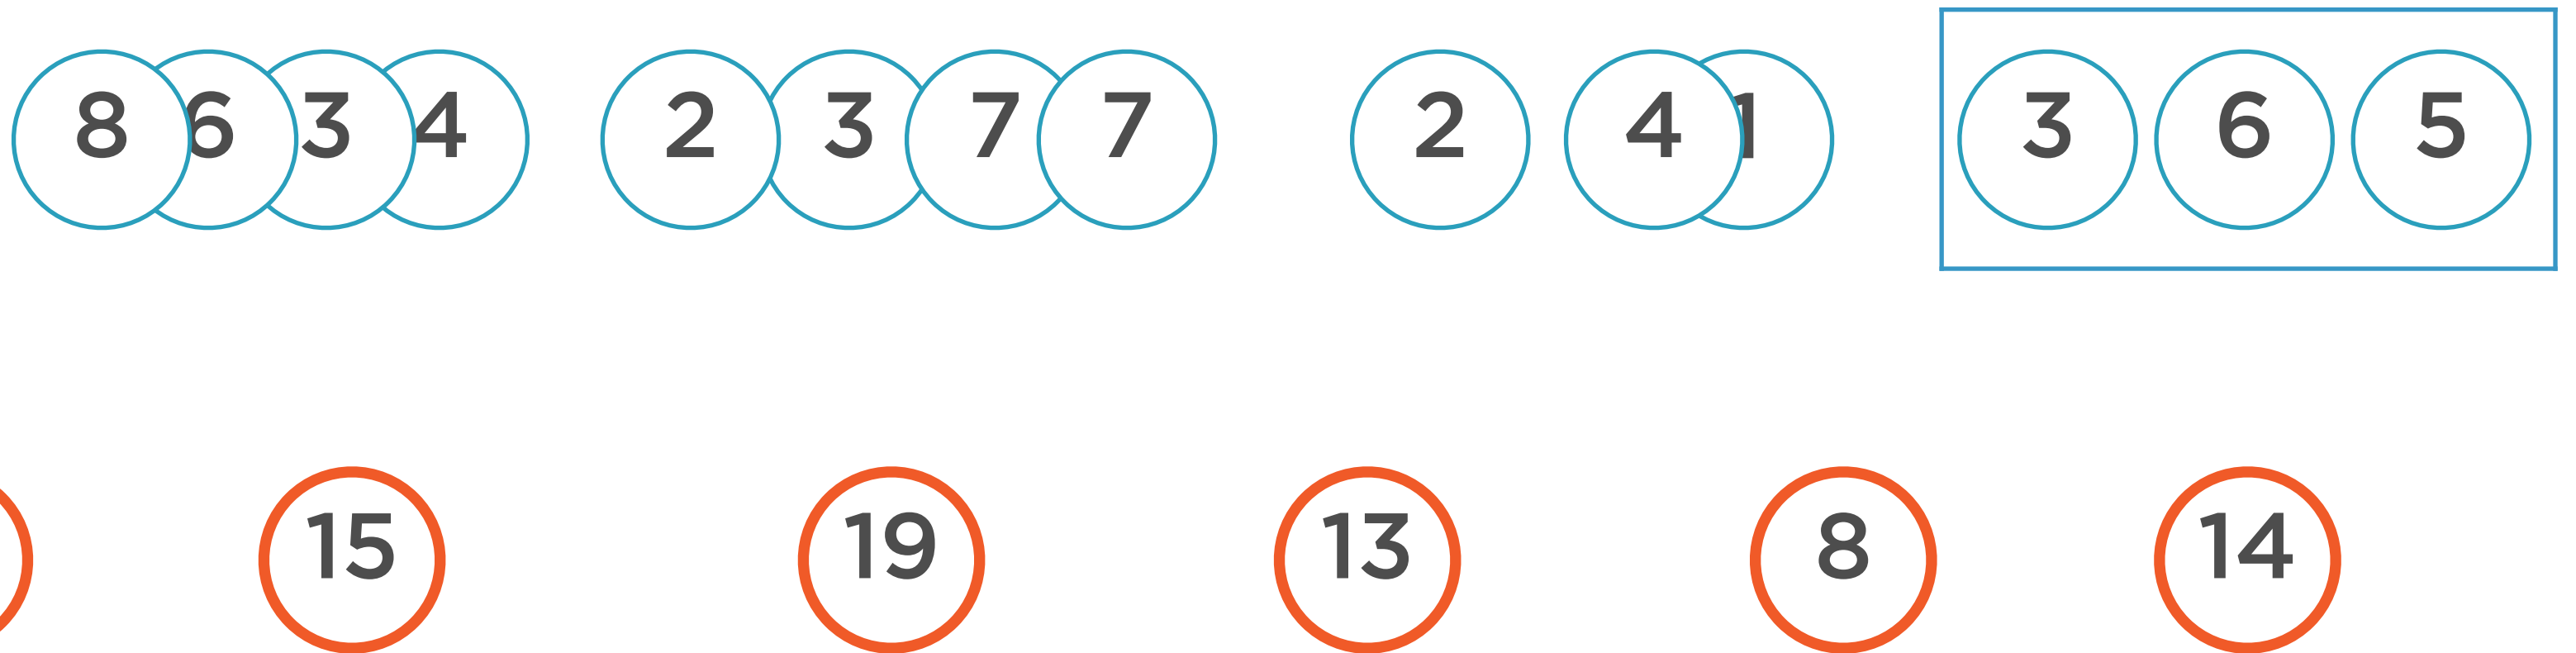


Fixed window size

**Overlapping** time - sliding interval


Number of entities differ within a window

# Sliding Window



Apply the `sum()` operation on each window

# Global Window




**All data in the stream in one window**

# Event Time and Processing Time

---

# Time-based Windows

Tumbling and sliding windows consider entities in a fixed interval of **time**



# Time-based Windows

Tumbling and sliding windows consider entities in a fixed interval of **time**

There are **different notions of time** that can apply to entities in a stream


# Time



**Event Time**



**Ingestion Time**



**Processing Time**

# Event Time



The time at which the event occurred at its **original** source

- Mobile phone, sensor, website

Usually **embedded** within records

Gives correct results in case of out of order or late events



# Ingestion Time




The time at which the event **enters the system** via a source

Timestamp given by system  
chronologically after the event time

Cannot handle out of order events

# Processing Time




The system time of the machine  
**processing** entities

Chronologically after event time and  
ingestion time


Non-deterministic, depends on when  
data arrives, how long operations take

Simple, no coordination between  
streams and processors


# Event Time vs. Processing Time




# Event Time vs. Processing Time



# Event Time vs. Processing Time



# Event Time vs. Processing Time



# Demo


**Exploring global windows, tumbling  
(fixed) windows, and sliding windows**

# Watermarks and Late Data

---



# How Late Is Late?



**Class At 9 am**

Class starts when  
clock strikes 9



**Is 9:01 Late?**

Realistically, at least  
some folks are going  
to be a minute late



**Is 10:10 late?**

A student is an hour  
late - allow in or send  
back?

# How Late Is Late?



**The professor “knows” what lateness is reasonable**

**Students entering within this reasonable lateness are late but OK**

**Students entering after this reasonable lateness are too late**

**“Allowed Lateness”**

# How Late Is Late?



## Dealing with excessive lateness

### A student is too late

- Option 1: Send back home
- Option 2: Allow in, continue class
- Option 3: Allow in, restart class(!)

# Watermarks and Late Data

The system “knows”  
what lateness is  
reasonable

Data entering  
within this  
reasonable lateness  
is late but OK

Data entering after  
this reasonable  
lateness is too late

# Watermarks and Late Data

## Watermark

Threshold of allowed  
lateness (event time)

Data entering  
within this  
reasonable lateness  
is late but OK

Data entering after  
this reasonable  
lateness is too late

# Watermarks and Late Data

## Watermark

Threshold of allowed  
lateness (event time)

## Late Data

Data within watermark is  
aggregated

Data entering after  
this reasonable  
lateness is too late

# Watermarks and Late Data

## Watermark

Threshold of allowed  
lateness (event time)

## Late Data

Data within watermark is  
aggregated

## Dropped Data

Data outside watermark is  
dropped

```
windowedCounts = words.groupBy(  
    window(words.timestamp, "10 minutes", "5 minutes"),  
    words.word  
) .count()
```

---

## Simple Group-by Without Watermark

**Count words in each sliding window of width 10 minutes, sliding by 5 minutes**



```
windowedCounts = words \  
    .withWatermark("timestamp", "12 minutes") \  
    .groupBy(  
        window(words.timestamp, "10 minutes", "5 minutes"),  
        words.word) \  
    .count()
```

---

## Simple Group-by With Watermark

**We define the watermark i.e. lateness threshold to be 12 minutes**

```
windowedCounts = words \  
    .withWatermark("timestamp", "12 minutes") \  
    .groupBy(  
        window(words.timestamp, "10 minutes", "5 minutes"),  
        words.word) \  
    .count()
```

---

## Simple Group-by With Watermark

**Now window triggering will be delayed by 12 minutes**

# Watermark



**System generated or user specified**

**If, say network speed drops, watermark can become more lenient**

**$\text{Lateness} = \text{Processing Time} - \text{Event time}$**

# Watermarks and Output Modes



**Append mode:** Window not triggered at all until watermark elapses

- No partial updates

**Update mode:** Window will trigger even before watermark elapses

- Engine will keep partial counts

**Complete mode:** Can not be used with watermarks

# Watermarks and Output Modes




**No complete-mode queries**

**Aggregation must be event-time, or event-time window**

**.withWatermark must be called on same timestamp column as aggregate**

**.withWatermark must be called before the aggregation**

# One-way Guarantee



All data **before** watermark will definitely not be dropped

All data **after** watermark may or may not be dropped

More delayed the data, less likely the engine is to process the data

Demo

**Using watermarks to allow late data**

# Demo

**Using UDFs (user-defined functions) to transform data**



# Summary

**Windowing operations on streams**

**Sliding and tumbling windows**

**Event time vs. processing time**

**Aggregation operations on windows**

**UDFs on streaming data**

**Up Next:**

Working with Streaming Joins

---